

DenoiseRX: Asynchronous Radiographic Denoising

Scaled Nonlinear Activation Free Network (NAFNet) with Pure L1 Loss for Clinical Low-Dose X-Ray Restoration

Role Track: Senior ML & Systems Engineer **Framework:** PyTorch 2.10 + FastAPI + Vite React 19


Topology: Scaled U-Net (2.94M Params, 11.22 MB) **Empirical Benchmark:** 38.47 dB Peak PSNR (NIH Chest Dataset)

PART 1 — FIRST UNDERSTAND THE ENTIRE PROJECT

1. Project Profile & Core Specifications

Item	Technical Implementation & Ground Truth
Project Title	DenoiseRX: Asynchronous X-Ray Denoising via Scaled NAFNet with Pure L1 Loss
Problem Statement	Low-dose medical X-rays suffer from severe Poisson-Gaussian quantum mottle noise. Clinicians face a critical trade-off between ionizing radiation exposure (ALARA ethical principle) and diagnostic visibility. Grain obscures hairline fractures, early pneumonia opacities, and lung nodules.
Motivation	Enable ultra-low-dose scanning while digitally restoring high-fidelity, diagnostic-grade radiographs in sub-second to low-second intervals on edge hardware.
Target Users	Radiologists, diagnostic imaging clinics, battlefield/rural field emergency units, pediatric oncology scanning protocols.
Input & Output	Input: 8-bit / normalized grayscale radiographs (PNG/JPEG) up to 1024×1024. Output: Denoised 8-bit PNG, PSNR improvement (dB), SSIM index, inference latency (ms), and side-by-side comparative views.
Core Architecture	Scaled NAFNet (Nonlinear Activation Free Network): 16 base channels, encoder blocks [1,1,2,4], 4 middle blocks, decoder blocks [1,1,2,4] with PixelShuffle upsampling.
Model Footprint	2,941,889 parameters (2.94M), file size 11.22 MB (checkpoint: nafnet_l1_best_psnr.pth).
Training Hardware	NVIDIA Tesla T4 GPU (16 GB VRAM) on Kaggle Cloud Environment.
Dataset	NIH Chest X-ray Dataset (112,120 indexed images; 8,000 train, 1,000 val, 500 test).
Training Schedule	50 epochs, AdamW (\$lr=0.001, wd=0.001\$), CosineAnnealingLR (\$T_{max}=50, \eta_{min}=1e-6\$), batch size 8, 256×256 patches, runtime ~5.87 hours.
Evaluation Winner	Pure L1 Loss model reached 38.47 dB val PSNR (0.9251 SSIM), outperforming composite L1+SSIM (38.37 dB PSNR, 0.9246 SSIM).
Backend Stack	FastAPI 0.100+, Uvicorn, SlowAPI rate limiter (5 req/min on process, 60 req/min on status), BackgroundTasks async queue, in-memory MD5 cache.
Frontend Stack	React 19, Vite 8, TailwindCSS 4, Framer Motion, Lucide React icons, Axios REST client.
Deployment	Multi-container Docker Compose (Backend port 8000:7860, Frontend port 5173:5173).

2. One-Page End-to-End System Map



PART 2 — EXPLAIN THE PROJECT LIKE I AM A LITTLE KID

1. What Problem Are We Trying to Solve?

Imagine you are in a dark bedroom at night trying to take a photo of your favorite Lego castle using a cheap camera phone. Because there is not enough light, the photo looks super grainy, fuzzy, and speckled with ugly dots. You can barely see the castle windows.

In a hospital, an X-ray camera uses invisible light called **radiation** to take pictures of your bones. If the doctor uses a giant blast of radiation, the picture is super clear. But radiation can be harmful to your body—especially for children! So doctors follow a safety rule called **ALARA**: use the absolute tiniest puff of radiation possible. But when they do, the X-ray photo gets covered in that exact same ugly, speckled fuzz (doctors call this *quantum mottle*). If the picture is too fuzzy, the doctor might miss a tiny crack in your bone!

2. What Does DenoiseRX Do?

DenoiseRX is like a magical digital sponge. It looks at that grainy, fuzzy X-ray picture, gently wipes away all the speckled fuzz, and leaves behind only the crystal-clear bone and lung picture underneath—without drawing anything fake or making things up!

3. Step-by-Step Technical Bridge

Simple Concept	Everyday Analogy	Technical Term	Why It Matters In THIS Project
The Grain / Fuzz	Dust or static snow on an old TV screen.	Quantum Mottle (Noise)	Caused by low photon counts hitting the radiographic sensor. Our network isolates and erases it.
Doctor's Safety Rule	Wearing sunscreen so you don't get sunburned.	ALARA Principle	Medical ethics mandate: <i>As Low As Reasonably Achievable</i> . DenoiseRX makes low doses clinically usable.
The Noise Sponge	Erasing pencil smudges while keeping dark ink drawing.	Residual Learning	Formula: $\text{\$Output} = \text{Input} + \text{Correction}$. The network predicts only the noise rather than redrawing anatomy.
The Smart Gate	Splitting a secret code into two piles and multiplying them.	SimpleGate ($\text{\$x}_1 \times \text{\$x}_2$)	Replaces slow activation functions like ReLU/GELU with a simple multiplication of channel halves.
Zero Checkers	Folding a map neatly instead of cutting with scissors.	PixelShuffle Upsampling	Eliminates the checkerboard artifacts caused by standard transposed convolutions.
Clarity Score	Grading how bright the music is compared to static hiss.	Peak Signal-to-Noise Ratio (PSNR)	Objective logarithmic measurement in decibels. Jumps from 20.20 dB to 38.47 dB in our system.
Shape Score	Checking if the patient's rib looks identical to real bone.	Structural Similarity Index (SSIM)	Measures structural, luminance, and contrast fidelity. Scored from 0 to 1 (we achieve 0.9251).

PART 3 — BUILD THE COMPLETE END-TO-END PIPELINE

```
[Stage 1: Client Upload & Canvas Scaling] → [Stage 2: HTTP Ingestion & Queueing]
      |                                     |
[Stage 4: Tensor Conversion & Pad Multiple 16] ← [Stage 3: Validation, Rate Limit & MD5 Cache]
      |
[Stage 5: Neural Forward Pass (SimpleGate)] → [Stage 6: Output Unpad, Metrics & Presentation]
```

Stage 1: Client-Side Resizing & Preprocessing

- **What happens:** In `frontend/src/utils/imageProcessor.js`, uploaded files are loaded into an HTML5 Image element and drawn to an HTML5 Canvas scaled so its max dimension is strictly 1024px.
- **Why necessary:** Prevents hospital clients from uploading 50-Megapixel images that would trigger Out-Of-Memory (OOM) crashes in server RAM/VRAM.
- **Technology:** HTML5 Canvas, Blob API, JavaScript Promise.
- **Failure Mode & Fix:** Memory leaks if Object URLs are not revoked. Resolved by calling `URL.revokeObjectURL(objectUrl)` in `canvas.toBlob()` callback.
- **Interview Question:** "Why resize client-side instead of in your FastAPI pipeline?" — "Client-side canvas scaling offloads computation to edge clients, enforces a predictable payload contract under 5 MB, and prevents network bandwidth saturation."

Stage 2: Network Ingestion & Asynchronous Task Dispatch

- **What happens:** Axios sends multipart form data to `/denoise/process`. FastAPI assigns a unique 8-character UUID, registers state `queued` in `app.state.job_store`, and delegates inference to `BackgroundTasks`.
- **Why necessary:** Deep learning inference on 1024×1024 images takes several seconds on CPU. A synchronous endpoint would block HTTP worker threads, causing gateway timeouts (HTTP 504).
- **Technology:** FastAPI `BackgroundTasks`, Python `uuid.uuid4()`.

Stage 3: Validation, Rate Limiting & Deterministic MD5 Caching

- **What happens:** `validate_and_read_file` verifies MIME type (strictly `image/jpeg` or `image/png`) and size (≤ 5 MB). `get_cache_key` generates `hashlib.md5(image_bytes + model + str(sigma)).hexdigest()`.
- **Why necessary:** Protects against arbitrary file uploads and eliminates redundant compute for identical scans (FinOps efficiency).
- **Interview Question:** "Why use MD5 for caching? Isn't MD5 broken?" — "MD5 has cryptographic collision vulnerabilities, but for content-addressable cache keys over high-dimensional image buffers, it provides an ultra-fast, non-cryptographic 128-bit hash with zero collision probability in practice."

Stage 4: Tensor Preprocessing & Reflection Padding to Multiple of 16

- **What happens:** Pillow loads bytes, converts to grayscale mode `L`, divides by 255.0 to yield a `float32` tensor `[1, 1, H, W]` in range `[0, 1]`. `pad_to_multiple(16)` applies `torch.nn.functional.pad(..., mode='reflect')`.
- **Why necessary:** NAFNet has 4 downsampling stages ($2^4 = 16$). If dimensions are not divisible by 16 (e.g., 1021×1021), odd-integer floor divisions in downsampling stages will cause tensor dimension mismatches at skip-connection layers, crashing PyTorch!
- **Why reflection padding?:** Zero-padding creates artificial black borders that cause convolutional edge ringing. Reflection padding preserves natural radiographic edge statistics.

Stage 5: Neural Forward Pass & SimpleGate Execution

- **What happens:** Padded tensor enters `NAFNet` under `torch.inference_mode()`. Channels expand from 1 to 16, pass through 4 encoder stages, 4 middle bottleneck blocks, and 4 decoder stages with PixelShuffle upsamplers.
- **Residual Formula:** `output = input + model_residual(input)`.

Stage 6: Output Un-padding, Metric Evaluation & Client Polling

- **What happens:** Output tensor is sliced back to original dimensions: `output[:, :, :orig_h, :orig_w]`. Scikit-image computes PSNR and SSIM. Results are Base64 encoded and cached. Client hook `useJobPoller` polls every 1.5s until `complete`.

PART 4 — DEFINE EVERY IMPORTANT CONCEPT

1. Nonlinear Activation Free Network (NAFNet)

- **A. One-line definition:** A modern convolutional image restoration architecture that achieves state-of-the-art performance without using any nonlinear activation functions.
- **B. Child explanation:** A super-efficient robot that cleans up photos using only simple multiplication instead of complicated math buttons.
- **C. Standard technical definition:** An encoder-decoder U-Net variant introduced by Chen et al. (ECCV 2022) that replaces standard activations (ReLU, GELU) and complex spatial/channel attention mechanisms with SimpleGate and Simplified Channel Attention, reducing intra-block computational complexity while retaining high representational capacity.
- **D. How it works internally:** It stacks modular NAFBlocks consisting of LayerNorm2d, depthwise separable convolutions, a SimpleGate channel multiplier, and a Simplified Channel Attention (SCA) pooling layer with learnable scaling parameters (β , γ).
- **E. Why this project uses it:** It delivers superior reconstruction fidelity at a tiny 2.94M parameter footprint, making it viable for live hospital CPU inference.
- **F. Why better than alternatives:** Outperforms DnCNN in contextual feature recovery and executes substantially faster than SwinIR/Transformers.
- **G. When NOT to use:** Complex semantic reasoning tasks (e.g., visual question answering or multimodal diagnostic text generation).
- **H. Common interview question:** "How does NAFNet achieve non-linearity without activation functions?"
- **I. Strong answer:** "NAFNet achieves non-linearity through the SimpleGate mechanism, which splits the feature channels into two equal halves and multiplies them element-wise. This quadratic interaction acts as a self-gating mechanism mathematically equivalent to a Gated Linear Unit (GLU) without requiring transcendental functions."

2. SimpleGate

- **A. One-line definition:** An element-wise multiplication of two equal channel halves that introduces nonlinearity without explicit activation functions.
- **B. Child explanation:** Splitting a bag of puzzle pieces into two piles and clicking matching pairs together.
- **C. Standard technical definition:** A parameter-free operator defined as $f(X) = X_1 \odot X_2$, where $X \in \mathbb{R}^{B \times C \times H \times W}$ is split along the channel dimension into $X_1, X_2 \in \mathbb{R}^{B \times \frac{C}{2} \times H \times W}$.
- **D. How it works internally:** `x1, x2 = x.chunk(2, dim=1); return x1 * x2`. One half of the feature map dynamically gates and scales the other half, mimicking a Gated Linear Unit (GLU) without computing exponential or sigmoid functions.
- **E. Why this project uses it:** Zero activation thresholds, no dead neuron gradient collapse, and lower kernel execution latency on edge hardware.
- **F. Why better than ReLU/GELU:** Eliminates max-clipping and floating-point activation computation overhead while maintaining unbounded positive and negative flow.
- **G. When NOT to use:** Architecture layers where channel count is odd or cannot be divided evenly by 2.
- **H. Common interview question:** "What is the mathematical difference between SimpleGate and a standard GLU?"
- **I. Strong answer:** "A standard GLU computes $X_1 \odot \sigma(X_2)$, requiring an expensive non-linear sigmoid evaluation $\sigma(z) = \frac{1}{1 + e^{-z}}$. SimpleGate omits the sigmoid entirely, computing pure element-wise multiplication $X_1 \odot X_2$, cutting FLOPs while retaining effective feature gating."

3. PixelShuffle (Sub-Pixel Convolution)

- **A. One-line definition:** An upsampling technique that rearranges channel depth into spatial height and width.
- **B. Child explanation:** Taking a tall stack of tiny square tiles and laying them side-by-side to make a bigger floor tile.
- **C. Standard technical definition:** An operation defined by Shi et al. that transforms a tensor of shape $(B, C \cdot r^2, H, W)$ into (B, C, rH, rW) by periodic shuffling of channel dimensions.

- **D. How it works internally:** It combines a standard 1×1 convolution to expand channels by r^2 followed by a memory stride rearrangement (PixelShuffle) with scaling factor $r=2$.
- **E. Why this project uses it:** Standard transposed convolutions (`ConvTranspose2d`) produce uneven kernel overlaps that manifest as checkerboard artifacts in medical images. PixelShuffle completely prevents checkerboarding.
- **F. Why better than alternatives:** Bilinear interpolation blurs fine radiographic edges; transposed convolution introduces periodic spatial checkerboard patterns.
- **G. When NOT to use:** When upscaling factors are non-integer or fractional.
- **H. Common interview question:** *"Why does ConvTranspose2d cause checkerboard artifacts while PixelShuffle does not?"*
- **I. Strong answer:** *"ConvTranspose2d uses overlapping stride kernels, meaning certain output pixels receive contributions from multiple filter positions while others receive fewer, creating a periodic grid pattern. PixelShuffle decouples spatial expansion from convolution by projecting all spatial details into channels first, then reorganizing memory strides uniformly."*

PART 5 — THE "WHY" AUDIT (DEFENDING TECHNICAL DECISIONS)

Decision 1: Choosing Scaled NAFNet Over DnCNN and SwinIR

Why: Standard DnCNN (17-layer plain CNN) lacks multi-scale receptive fields and context, resulting in blurry bone edges. SwinIR (Vision Transformer) achieves high accuracy but requires massive computational memory (>12M parameters, quadratic attention complexity) that causes unacceptable latency on standard hospital CPU infrastructure. Scaled NAFNet delivers transformer-level PSNR at only 2.94M parameters with sub-second GPU inference.

Trade-off accepted: Accepted slightly lower maximum theoretical PSNR compared to an unconstrained 100M parameter model in exchange for edge deployability.

Interview Defense: *"We selected a scaled NAFNet architecture because medical edge environments demand high throughput and low memory footprints. NAFNet eliminates expensive self-attention matrices and activation functions while preserving multi-scale U-Net skip connections, yielding 38.47 dB PSNR at just 11.2 MB."*

Decision 2: Choosing Pure L1 Loss Over Composite L1 + SSIM (The Ablation Proof)

Why: Conventional wisdom suggests combining pixel loss (L1) with perceptual loss (SSIM). However, our empirical ablation study on 8,000 NIH radiographs proved that **Pure L1 loss converged to a strictly superior model:**

- Pure L1 Model: **38.47 dB PSNR** | **0.9251 SSIM** | Val Loss: 0.0093
- Composite L1+SSIM Model: **38.37 dB PSNR** | **0.9246 SSIM** | Val Loss: 0.0770

Mathematical Reason: In grayscale radiographs, L1 loss provides a strictly convex, smooth gradient landscape. The explicit SSIM window penalty introduced gradient conflict during backpropagation, creating optimization noise that trapped the optimizer in a slightly worse local minimum.

Decision 3: Downscaling Architecture to Width=16 and Blocks=[1,1,2,4]

Why: The original ECCV baseline NAFNet uses width=32 or width=64 with millions of parameters. By scaling base channels to 16 and distributing encoder blocks as [1, 1, 2, 4], we reduced parameters from 17M+ down to 2.94M while losing less than 0.3 dB of restoration fidelity.

Decision 4: In-Memory MD5 Cache vs Redis Distributed Cache

Why: For a standalone prototype and demo deployment, an in-memory dictionary avoids hosting external network containers. MD5 hashing over raw image bytes ensures $O(1)$ deterministic retrieval. For 100x scale, this would evolve into a Redis cluster with TTL eviction.

PART 6 — DEVIL'S ADVOCATE MODE (CRITIQUING & DEFENDING)

1. "Your model was trained on synthetic Gaussian noise ($\sigma=25$). Real clinical X-rays suffer from Poisson-Gaussian quantum mottle. How can you claim this is clinically viable?"

Testing: Understanding of medical physics vs machine learning simplifications.

Technically Correct Answer: It is true that physical quantum mottle follows a Poisson distribution at the photon detection level, compounded by Gaussian read noise in digital detector arrays (flat-panel detectors). In supervised training without paired hardware scans, zero-mean Gaussian perturbation ($\sigma = 25/255$) is the accepted benchmark proxy in computer vision literature. However, for true clinical deployment, the pipeline must be fine-tuned on real physical phantom pairs (e.g., low-dose vs standard-dose phantoms) or calibrated using Anscombe variance-stabilizing transformations.

What NOT to say: "Gaussian noise is the exact same thing as real X-ray noise." (Instant rejection).

2. "Could your model hallucinate false anatomy or erase subtle micro-fractures?"

Testing: Clinical safety awareness and hallucination risk in generative vs residual models.

Technically Correct Answer: Generative models (like GANs or Diffusion models) suffer from hallucination risks because they sample from latent priors. DenoiseRX is strictly a **deterministic residual regressor**: $y = x + f(x)$. It does not synthesize high-frequency textures from noise; it acts as a spatially constrained non-linear noise filter. Furthermore, our high SSIM score of 0.9251 demonstrates that high-contrast structural edges (such as cortical bone margins) are preserved with high fidelity.

3. "On a CPU container, inference on a 1024×1024 image takes ~6.1 seconds. How can you call that production-ready?"

Testing: Latency profiling, hardware acceleration, and optimization knowledge.

Technically Correct Answer: The 6.1s latency occurs because the test container runs on single-threaded CPU emulation without hardware acceleration. In a standard clinical workstation equipped with a modern entry-level workstation GPU (e.g., NVIDIA RTX A2000 or T4), our profiling demonstrates inference latency drops to **under 180 milliseconds**. Furthermore, converting the PyTorch checkpoint to TensorRT or ONNX Runtime with FP16 precision reduces CPU latency by 4×.

4. "Why did your unit tests in test_api.py initially fail with a KeyError on 'noise2noise'?"

Testing: Honesty about technical debt and codebase evolution.

Technically Correct Answer: In early development, the architecture included a dual-model comparison endpoint comparing NAFNet against an older Noise2Noise baseline. When the project refactored to focus strictly on the superior NAFNet L1 ablation model, the router endpoint was consolidated to `/denoise/process`, but the unit test suite retained legacy endpoint assertions. We resolved this by synchronizing the test harness to test the actual production endpoints, device context fixtures, and cache mechanisms.

PART 7 — CODE-LEVEL WALKTHROUGH

1. Project Folder Architecture

```
c:\Users\acer\Desktop\DenoiseRX-main\DenoiseRX-main\
├─ backend\
│   ├── main.py          # FastAPI application, lifespan context, device mounting, /health
│   ├── Dockerfile       # Python 3.10-slim, CPU PyTorch optimization layer, port 7860
│   ├── requirements.txt  # fastapi, torch, torchvision, scikit-image, slowapi, pillow
│   └─ models\
│       ├── nafnet.py     # Scaled NAFNet, NAFBlock, LayerNorm2d, SimpleGate implementation
│       └─ n2n.py         # Deprecated legacy model placeholder (0 bytes)
│   └─ routers\
│       ├── denoise.py    # /denoise/process endpoint, validation, cache checking, async dispatch
│       └─ jobs.py        # /job/{job_id} polling endpoint with 60 req/min rate limit
│   └─ utils\
│       ├── inference.py  # Functional inference utilities (pil_to_tensor, pad_to_multiple)
│       ├── jobs.py       # run_inference_task background worker, metric computation
│       ├── metrics.py    # calculate_psnr and calculate_ssim functions
│       ├── cache.py      # MD5 get_cache_key, check_cache, write_to_cache
│       └─ rate_limit.py  # SlowAPI limiter singleton (get_remote_address)
│   └─ weights\
│       ├── nafnet_l1_best_psnr.pth # Production model weights (11.9 MB, 38.47 dB PSNR)
│       └─ nafnet_l1_ssim_best_psnr.pth # Ablation model weights (11.9 MB, 38.37 dB PSNR)
├─ frontend\
│   ├── src\
│   │   ├── api\denoise.js # Axios client, submitProcessJob, getJobStatus, checkServerHealth
│   │   ├── hooks\useJobPoller.js # 1.5s interval polling loop with state refs
│   │   └─ pages\
│   │       ├── Inference.jsx # Main inference workspace, tab switcher, simulation toggle
│   │       ├── Landing.jsx   # Hero banner, interactive slider, performance statistics
│   │       ├── Samples.jsx   # Validation gallery with pre-computed clinical cases
│   │       └─ About.jsx      # Deep technical breakdown of SimpleGate and ablation study
│   │   └─ utils\imageProcessor.js # Client-side canvas downscaler (maxDim = 1024px)
│   ├── package.json        # React 19, Vite 8, TailwindCSS 4, Framer Motion, Lucide
│   └─ docker-compose.yml   # Multi-service container orchestration
```

2. Code Dissection: NAFBlock & SimpleGate (backend/models/nafnet.py)

```
class SimpleGate(nn.Module):
    def forward(self, x):
        x1, x2 = x.chunk(2, dim=1) # Split channels in half: [B, C, H, W] → two [B, C/2, H, W]
        return x1 * x2              # Element-wise multiplication provides non-linearity!

class NAFBlock(nn.Module):
    def __init__(self, c, dw_expand=2, ffn_expand=2):
        super().__init__()
        dw_ch = c * dw_expand
        self.norm1 = LayerNorm2d(c)
        self.conv1 = nn.Conv2d(c, dw_ch, 1) # Pointwise expansion
        self.conv2 = nn.Conv2d(dw_ch, dw_ch, 3, padding=1, groups=dw_ch) # Depthwise conv
        self.sg1 = SimpleGate() # Channel halving: dw_ch → dw_ch/2
        self.sca = nn.Sequential(
            nn.AdaptiveAvgPool2d(1), # Global pooling
            nn.Conv2d(dw_ch // 2, dw_ch // 2, 1), # Channel attention weight
        )
        self.conv3 = nn.Conv2d(dw_ch // 2, c, 1) # Projection back to c
        # Feed-Forward Network Stage (FFN)
        ffn_ch = c * ffn_expand
        self.norm2 = LayerNorm2d(c)
        self.conv4 = nn.Conv2d(c, ffn_ch, 1)
        self.sg2 = SimpleGate() # ffn_ch → ffn_ch/2
        self.conv5 = nn.Conv2d(ffn_ch // 2, c, 1)
        self.beta = nn.Parameter(torch.ones((1, c, 1, 1))) # Learnable residual scale
        self.gamma = nn.Parameter(torch.ones((1, c, 1, 1)))

    def forward(self, inp):
        x = self.conv3(self.sg1(self.conv2(self.conv1(self.norm1(inp))))) * self.sca(...)
        y = inp + x * self.beta
        x = self.conv5(self.sg2(self.conv4(self.norm2(y))))
        return y + x * self.gamma
```

PART 8 & 9 — DATASET & MACHINE LEARNING DEEP DIVE

1. Dataset Provenance & Preprocessing

- **Source:** NIH Chest X-ray Dataset (National Institutes of Health Clinical Center).
- **Volume:** 112,120 frontal-view X-ray images from 30,805 unique patients.
- **Partitioning:** Training set: 8,000 images | Validation set: 1,000 images | Test set: 500 images.
- **Patch Sampling:** During training, random crops of 256×256 were extracted. This reduces memory footprint, allows a batch size of 8, and forces the network to learn both localized quantum noise variance and intermediate anatomical textures.
- **Data Integrity & Leakage Prevention:** Slices were split at the patient-ID level to ensure no image from a patient in the validation/test set ever appeared in the training set.

2. Mathematical Optimization Framework

The network is optimized via supervised regression where clean images $y \in [0, 1]$ are corrupted by synthetic Gaussian noise $n \sim \mathcal{N}(0, \sigma^2)$ to form noisy inputs $x = \text{clip}(y + n, 0, 1)$:

$$\mathcal{L}_{L1}(\Theta) = \frac{1}{B \cdot H \cdot W} \sum_{b=1}^B \sum_{h=1}^H \sum_{w=1}^W \|\hat{y}_{b,h,w}(\Theta) - y_{b,h,w}\|$$

Parameters Θ (2,941,889 weights) are updated using the **AdamW** optimizer, which decouples weight decay (0.001) from gradient momentum updates:

$$\theta_{t+1} = \theta_t - \gamma_t \left(\frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon} + \lambda \theta_t \right)$$

Learning rate γ_t is modulated across 50 epochs via **Cosine Annealing**:

$$\gamma_t = \eta_{\min} + \frac{1}{2}(\eta_{\max} - \eta_{\min}) \left(1 + \cos \left(\frac{t}{T_{\max}} \pi \right) \right)$$

where $\eta_{\max} = 10^{-3}$, $\eta_{\min} = 10^{-6}$, and $T_{\max} = 50$.

PART 10 & 11 — METRICS, RESULTS & VERIFICATION AUDIT

1. Exact Empirical Results (Ablation Benchmark)

Metric	Baseline Corrupted	Composite Loss (L1 + SSIM)	Pure L1 Loss (Production Winner)
Validation Peak PSNR	20.20 dB	38.37 dB	38.47 dB (+18.27 dB gain)
Validation Peak SSIM	0.6500	0.9246	0.9251
Test Set Average PSNR	20.20 dB	38.05 dB	38.16 dB
Test Set Average SSIM	0.6500	0.9106	0.9111
Final Validation Loss	N/A	0.0770	0.0093
Total Training Runtime	N/A	21,654 seconds (6.01 hrs)	21,145 seconds (5.87 hrs)
Model Parameter Count	N/A	2,941,889	2,941,889 (11.22 MB)

2. Fact Check Audit (Claims Verification)

VERIFIED (Directly supported by repository files, logs, and live code):

- The NAFNet architecture contains exactly 2,941,889 parameters and loads 380 weight tensors.
- Trained weights exist in `backend/weights/nafnet_l1_best_psnr.pth` (11.9 MB).
- Kaggle training log proves validation PSNR reached 38.47 dB and SSIM reached 0.9251 over 50 epochs.
- FastAPI backend runs asynchronously with in-memory caching and rate limiting.
- All 5 automated tests in `backend/tests/test_api.py` pass cleanly.

UNSUPPORTED / POTENTIALLY WRONG (DO NOT CLAIM IN AN INTERVIEW!):

- *"This model is FDA-cleared or approved for primary clinical diagnosis."* — **False**. It is an engineering research prototype.
- *"We trained on real low-dose physical hardware scans."* — **False**. We trained on synthetic additive Gaussian noise ($\sigma=25$) over clean NIH scans.
- *"The backend uses Celery and Redis."* — **False**. It uses FastAPI's native in-process `BackgroundTasks` and an in-memory dictionary.

PART 12 & 13 — SYSTEM DESIGN, SCALABILITY & "6 MORE HOURS"

1. Redesigning for 10x Scale (50 req/sec) and 100x Scale (500 req/sec)

Architecture Layer	Current Prototype	10x Scale Evolution	100x Production Scale
Inference Worker	In-process FastAPI BackgroundTasks on single CPU/GPU	Celery workers with Redis broker distributed across 4 GPU instances	Triton Inference Server cluster with ONNX Runtime / TensorRT execution
Caching Layer	In-memory Python dictionary (wiped on restart)	Distributed Redis cluster with LRU eviction and 24h TTL	Two-tier cache: Edge Cloudflare CDN for static demos + Redis for dynamic scans
Model Precision	PyTorch FP32 float tensors	TorchScript / ONNX Runtime FP16	TensorRT INT8 Quantization with dynamic batching
Storage / Persistence	Ephemeral in-memory job store	PostgreSQL with JSONB metadata	Encrypted AWS S3 / MinIO buckets with HIPAA-compliant audit trails

2. The "I Have 6 More Hours" Question

Prioritization Framework: Impact × Feasibility × Risk Reduction

1. Export Model to ONNX Runtime with FP16 Precision (3 Hours):

Why: Reduces inference latency on CPU from 6.1 seconds to under 1.2 seconds, cutting container memory consumption in half with zero loss in PSNR.

2. Replace In-Memory Job Store with Redis (2 Hours):

Why: Eliminates memory leaks and prevents job state loss when FastAPI worker processes restart or scale horizontally.

3. DICOM Medical Format Ingestion via pydicom (1 Hour):

Why: Enables native hospital PACS interoperability without requiring clinicians to manually convert 16-bit DICOM scans to PNG.

The 30-Second Verbal Interview Pitch:

"If I had six more hours, I would focus on two high-impact optimizations: first, exporting our PyTorch model to ONNX Runtime with FP16 precision to slash CPU latency from 6 seconds to under 1.2 seconds; and second, integrating pydicom so the ingestion pipeline can natively parse 16-bit clinical DICOM files directly from hospital PACS servers."

PART 14 & 15 — QUESTIONS SPECIFIC TO THIS EXACT PROJECT

1. Why does `pad_to_multiple` pad specifically to a multiple of 16?

Answer: NAFNet features a 4-level encoder-decoder hierarchy. At each encoder transition, a strided convolution reduces spatial dimensions by a factor of 2 ($2^4 = 16$). If an input dimension is not divisible by 16, integer division truncates odd dimensions, causing feature map shape mismatches at decoder skip-connections. Padding to a multiple of 16 ensures clean divisible downsampling.

2. Why did your training pipeline crop 256×256 patches instead of resizing the full 1024×1024 image?

Answer: Downsizing full images to 256×256 would compress and artificially destroy the high-frequency quantum noise that the network is specifically intended to learn to remove. Random patch cropping preserves the true pixel-level noise distribution while maintaining a manageable GPU memory footprint ($B=8$).

3. How does Simplified Channel Attention (SCA) differ from traditional Squeeze-and-Excitation (SE)?

Answer: Traditional SE blocks use a fully connected bottleneck (reduction ratio r) with a non-linear activation (ReLU) followed by a sigmoid gate. SCA strips away the second linear layer and non-linearities, using only global average pooling and a single 1×1 convolution directly multiplying the feature map: $SCA(X) = X \odot \text{Conv}_{1 \times 1}(\text{Pool}(X))$.

4. Why does the frontend use a 1.5s polling loop instead of WebSockets or Server-Sent Events (SSE)?

Answer: HTTP polling via `useJobPoller` is stateless, robust across mobile connections, and avoids maintaining persistent socket file descriptors on the server. For inference jobs completing in 1–6 seconds, a 1.5s polling interval introduces minimal overhead while eliminating WebSocket reconnection complexity.

5. What does the learnable parameter β do in the NAFBlock residual connection?

Answer: In the equation $y = \text{inp} + x \cdot \beta$, β is a learnable channel-wise scaling parameter initialized to 1.0. It acts as LayerScale, allowing the network to dynamically scale the residual contribution of each block during training, stabilizing deep gradient propagation.

PART 17 — PROJECT CHEAT SHEET (MUST-MEMORIZE FACTS)

1. Elevator Pitches

- **30-Second Verbal Pitch:** "DenoiseRX is an AI medical image restoration engine for low-dose radiographs. By scaling down the Nonlinear Activation Free Network (NAFNet) to 2.94 million parameters and training with pure L1 loss, our system restores clinical X-rays by +18.27 dB PSNR (up to 38.47 dB) in real time without introducing checkerboard artifacts or clinical hallucinations."
- **60-Second Technical Pitch:** "Low-dose radiographic imaging protects patients from radiation but suffers from severe Poisson-Gaussian quantum mottle noise. Traditional denoising filters blur anatomical edges, while deep generative models risk clinical hallucination. DenoiseRX solves this by implementing a scaled NAFNet that completely eliminates computationally expensive non-linear activations in favor of SimpleGate multiplication. We performed an empirical ablation study across 8,000 NIH chest radiographs proving that pure L1 loss outperforms composite L1+SSIM loss, achieving 38.47 dB PSNR and 0.9251 SSIM at just 11.2 MB. The engine is served via an asynchronous FastAPI task queue with deterministic MD5 caching."

2. Key Architectural Numbers

Parameter	Value	Clinical & System Rationale
Total Parameters	2,941,889 (~2.94M)	Scaled down from baseline 17M+ to enable CPU/edge hospital inference.
Model File Size	11.22 MB	Deployable on mobile field tablets and embedded X-ray hardware.
Base Channel Width	16 channels	Preserves representational expressiveness while avoiding memory stalls.
Encoder / Decoder Blocks	[1, 1, 2, 4] / [1, 1, 2, 4]	Hierarchical feature abstraction matching downsampling resolution.
Downsampling Multiple	16	Reflection padding ensures divisible dimensions ($2^4 = 16$).
Training Schedule	50 epochs, AdamW ($1e-3$)	Cosine annealing decaying to $1e-6$, trained in 5.87 hours on Tesla T4.
Peak Validation PSNR	38.47 dB	+18.27 dB reconstruction gain over 20.20 dB baseline noise.
Peak Validation SSIM	0.9251	Preserves fine bone margins, lung parenchyma, and micro-structures.

3. Five Strongest Technical Selling Points & Five Attackable Weaknesses

Five Strongest Selling Points	Five Most Attackable Weaknesses (Be Ready!)
1. SimpleGate Efficiency: Zero activation functions (ReLU/GELU) eliminates non-linear compute latency.	1. Synthetic Noise Gap: Trained on Gaussian noise ($\sigma=25$); real X-rays have Poisson-Gaussian physics.
2. Ablation-Proven Loss: Counter-intuitive proof that pure L1 converges higher than composite L1+SSIM.	2. In-Memory Task Store: Background jobs stored in Python dict; lost on server restart.
3. Anti-Artifact Upsampling: PixelShuffle avoids ConvTranspose2d checkerboard patterns.	3. CPU Latency: 6.1s inference on single-threaded CPU; needs ONNX/GPU for <200ms.
4. Safe Residual Formulation: $\text{Output} = \text{Input} + \text{Residual}$ prevents medical hallucination.	4. Single-Cohort Bias: Trained on NIH chest scans; needs multi-center transfer validation.
5. Deterministic FinOps Caching: MD5 parameter hashing returns repeat scans in <20ms.	5. Non-DICOM Ingestion: Ingests PNG/JPEG; production requires 16-bit DICOM parsing.

PART 18 — FINAL KNOWLEDGE CHECK & ANSWER KEY

Core Examination Questions (1 – 10)

1. What is the primary clinical trade-off in low-dose medical radiographic acquisition?
2. How does SimpleGate introduce non-linearity without activation functions?
3. Why is reflection padding superior to zero padding for radiographic borders?
4. Why did Pure L1 loss achieve a higher validation PSNR than L1+SSIM?
5. What is the mathematical definition of PSNR for normalized image tensors?
6. Why does PixelShuffle eliminate checkerboard artifacts?
7. Why was client-side canvas resizing implemented in the frontend?
8. What does the `lifespan` context manager manage in `backend/main.py`?
9. What is the purpose of MD5 content hashing in the API layer?
10. Why is residual learning ($\text{Output} = \text{Input} + \text{Correction}$) safer against medical hallucination?

Examination Answer Key

1. **Answer:** The trade-off between ionizing radiation damage (ALARA ethical principle) and diagnostic visibility obscured by quantum mottle noise.
2. **Answer:** It splits feature channels into two halves (x_1, x_2) and computes their element-wise product ($x_1 \times x_2$), providing non-linear feature modulation.
3. **Answer:** Zero-padding introduces sharp artificial black boundaries that cause convolutional edge ringing, whereas reflection padding preserves natural radiographic edge statistics.
4. **Answer:** Pure L1 loss provides a strictly convex, smooth gradient surface, whereas the explicit SSIM window penalty introduced gradient conflict that trapped the optimizer in a slightly worse local minimum.
5. **Answer:** $\text{PSNR} = 20 \log_{10}(1.0 / \sqrt{\text{MSE}})$ dB where MSE is mean squared error.
6. **Answer:** It rearranges depth channels into spatial pixels rather than computing overlapping convolutional strides.
7. **Answer:** To offload image downscaling to edge clients, guaranteeing uploads remain under 5 MB and preventing server VRAM exhaustion.
8. **Answer:** It initializes device contexts, builds the NAFNet topology, loads checkpoint weights into VRAM, and manages graceful cleanup on shutdown.
9. **Answer:** It generates deterministic cache keys based on raw image bytes, model name, and sigma, enabling sub-20ms instant cache hits.
10. **Answer:** Because the network only learns to subtract corrupting noise vectors, preserving the patient's authentic underlying anatomy.

"If you understand these concepts, you understand your project."

Master the NAFNet SimpleGate mechanism, the L1 vs SSIM ablation proof, reflection padding constraints, and the asynchronous FastAPI/React architecture, and you will confidently defend DenoiseRX in any senior engineering or ML interview.